

# An introduction to R: **Algorithmics in R (continued)**

Noémie Becker, Benedikt Holtmann & Dirk Metzler <sup>1</sup>

nbecker@bio.lmu.de - holtmann@bio.lmu.de

Winter semester 2016-17

---

<sup>1</sup>Special thanks to: Prof. Dr. Martin Hutzenthaler and Dr. Sonja Grath for course development

- 1 Writing your own functions
- 2 `apply()` and `tapply()`
- 3 How to avoid slow R code

# Contents

- 1 Writing your own functions
- 2 `sapply()` and `tapply()`
- 3 How to avoid slow R code

# Basics

Syntax:

```
myfun <- function (arg1, arg2, . . .) { commands }
```

# Basics

Syntax:

```
myfun <- function (arg1, arg2, . . .) { commands }
```

Example:

We want to define a function that takes a DNA sequence as input and gives as output the GC content (proportion of G and C in the sequence).

# Basics

Syntax:

```
myfun <- function (arg1, arg2, . . .) { commands }
```

Example:

We want to define a function that takes a DNA sequence as input and gives as output the GC content (proportion of G and C in the sequence).

```
?gc # oops there is already a function named gc
```

# Basics

Syntax:

```
myfun <- function (arg1, arg2, . . .) { commands }
```

Example:

We want to define a function that takes a DNA sequence as input and gives as output the GC content (proportion of G and C in the sequence).

```
?gc # oops there is already a function named gc
```

```
?GC # ok this time
```

# Basics

Syntax:

```
myfun <- function (arg1, arg2, . . .) { commands }
```

Example:

We want to define a function that takes a DNA sequence as input and gives as output the GC content (proportion of G and C in the sequence).

```
?gc # oops there is already a function named gc
```

```
?GC # ok this time
```

We will use the function `gregexpr` for regular expressions.

```
?gregexpr
```

```
GC <- function(dna) {  
  gc.cont <- length(gregexpr("C|G",dna)[[1]])/nchar(dna)  
  return(gc.cont)  
}
```



# Basics

Syntax:

```
myfun <- function (arg1, arg2, . . .) { commands }
```

Example:

We want to define a function that takes a DNA sequence as input and gives as output the GC content (proportion of G and C in the sequence).

```
?gc # oops there is already a function named gc
```

```
?GC # ok this time
```

We will use the function `gregexpr` for regular expressions.

```
?gregexpr
```

```
GC <- function(dna) {  
  gc.cont <- length(gregexpr("C|G",dna)[[1]])/nchar(dna)  
  return(gc.cont)  
}
```

```
GC("AATTCGCTTA")
```

```
[1] 0.3
```

# Are we sure our function is correct?

# Are we sure our function is correct?

```
GC("AATTAAATTA")
```

# Are we sure our function is correct?

```
GC("AATTAAATTA")
```

```
[1] 0.1
```

What happened?

# Are we sure our function is correct?

```
GC("AATTAAATTA")
```

```
[1] 0.1
```

What happened?

A function should always be tested with several inputs.

# Better version of the function

```
GC <- function(dna) {  
  gc1 <- gregexpr("C|G",dna)[[1]]  
  if (length(gc1)>1){  
    gc.cont <- length(gc1)/nchar(dna)  
  } else {  
    if (gc1>0){  
      gc.cont <- 1/nchar(dna)  
    } else {  
      gc.cont <- 0  
    }  
  }  
  return(gc.cont)  
}
```

# Deal with wrong arguments

So far we assumed that the input was a chain of characters with only A, T, C and G.

What happens if we try another type of argument?

```
GC("23")
```

```
[1] 0
```

# Deal with wrong arguments

So far we assumed that the input was a chain of characters with only A, T, C and G.

What happens if we try another type of argument?

```
GC("23")
```

```
[1] 0
```

```
GC(TRUE)
```

```
[1] 0
```



# Deal with wrong arguments

So far we assumed that the input was a chain of characters with only A, T, C and G.

What happens if we try another type of argument?

```
GC("23")
```

```
[1] 0
```

```
GC(TRUE)
```

```
[1] 0
```

```
GC("notDNA")
```

```
[1] 0
```

# Deal with wrong arguments

So far we assumed that the input was a chain of characters with only A, T, C and G.

What happens if we try another type of argument?

```
GC("23")
```

```
[1] 0
```

```
GC(TRUE)
```

```
[1] 0
```

```
GC("notDNA")
```

```
[1] 0
```

```
GC("Cool")
```

```
[1] 0.25
```

# Deal with wrong arguments

So far we assumed that the input was a chain of characters with only A, T, C and G.

What happens if we try another type of argument?

```
GC("23")
```

```
[1] 0
```

```
GC(TRUE)
```

```
[1] 0
```

```
GC("notDNA")
```

```
[1] 0
```

```
GC("Cool")
```

```
[1] 0.25
```

How can we deal with this?

What do we want our function to output in these cases?

Find a solution collectively (answer below).

# Error and warning

There are two types of error messages in R:

- Error message stops execution and returns no value.
- Warning message continues execution.

# Error and warning

There are two types of error messages in R:

- Error message stops execution and returns no value.
- Warning message continues execution.

```
x <- sum("hello")
```

```
Error in sum("hello") : invalid 'type' (character) of argument
```

# Error and warning

There are two types of error messages in R:

- Error message stops execution and returns no value.
- Warning message continues execution.

```
x <- sum("hello")
```

```
Error in sum("hello") :  invalid 'type' (character) of argument
```

```
x <- mean("hello")
```

```
Warning message:
```

```
In mean.default("hello") :  argument is not numeric or logical:  
returning NA
```

# Error and warning

There are two types of error messages in R:

- Error message stops execution and returns no value.
- Warning message continues execution.

```
x <- sum("hello")
```

```
Error in sum("hello") : invalid 'type' (character) of argument
```

```
x <- mean("hello")
```

```
Warning message:
```

```
In mean.default("hello") : argument is not numeric or logical:  
returning NA
```

We can define such messages with the functions `stop()` and `warning()`.

In our example:

- Error when argument not character
- Warning if character argument not DNA.

# Deal with non character arguments

```
GC <- function(dna) {  
  if (!is.character(dna)){  
    stop("The argument must be of type character.")  
  }  
  gc1 <- gregexpr("C|G",dna)[[1]]  
  if (length(gc1)>1){  
    gc.cont <- length(gc1)/nchar(dna)  
  } else {  
    if (gc1>0){  
      gc.cont <- 1/nchar(dna)  
    } else {  
      gc.cont <- 0  
    }  
  }  
  return(gc.cont)  
}
```



# Deal with non DNA character

We define as non DNA any character different from A, C, T, G.  
If there is another character we compute the value but issue a warning.

# Deal with non DNA character

We define as non DNA any character different from A, C, T, G.  
If there is another character we compute the value but issue a warning.

We can use the function `grep` as follows:

```
grep("[^GCAT]", dna)
integer(0)
grep("[^GCAT]", "fATCG")
[1] 1
```

# Deal with non DNA character

```
GC <- function(dna) {  
  if (!is.character(dna)){  
    stop("The argument must be of type character.")  
  }  
  if (length(grep("[^GCAT]", dna))>0){  
    warning("The input contains characters other than A, C, T, G - value  
should not be trusted!")  
  }  
  gc1 <- gregexpr("C|G",dna)[[1]]  
  if (length(gc1)>1){  
    gc.cont <- length(gc1)/nchar(dna)  
  } else {  
    if (gc1>0){  
      gc.cont <- 1/nchar(dna)  
    } else {  
      gc.cont <- 0  
    }  
  }  
  return(gc.cont)  
}
```

# Giving several arguments to a function

Most R functions have several arguments. You can see them listed in the help page.

# Giving several arguments to a function

Most R functions have several arguments. You can see them listed in the help page.

A frequent argument in R functions is `na.rm` that removes NA values from vectors if it is set to `TRUE`.

```
mean(c(1,2,NA))
```

```
[1] NA
```

```
mean(c(1,2,NA), na.rm=TRUE)
```

```
[1] 1.5
```

We could give our function a second argument to output the AT content instead of GC.

# Giving several arguments to a function

```
GC <- function(dna,AT ) {  
  gc1 <- gregexpr("C|G",dna)[[1]]  
  if (length(gc1)>1){  
    gc.cont <- length(gc1)/nchar(dna)  
  } else {  
    if (gc1>0){  
      gc.cont <- 1/nchar(dna)  
    } else {  
      gc.cont <- 0  
    }  
  }  
  if (AT==TRUE){  
    return(1-gc.cont)  
  } else {  
    return(gc.cont)  
  }  
}
```

# Giving several arguments to a function

```
GC <- function(dna,AT ) {  
  gc1 <- gregexpr("C|G",dna)[[1]]  
  if (length(gc1)>1){  
    gc.cont <- length(gc1)/nchar(dna)  
  } else {  
    if (gc1>0){  
      gc.cont <- 1/nchar(dna)  
    } else {  
      gc.cont <- 0  
    }  
  }  
  if (AT==TRUE){  
    return(1-gc.cont)  
  } else {  
    return(gc.cont)  
  }  
}
```

Test:

```
GC(dna,AT=TRUE) [1] 0.7
```

# Giving a default value to an argument

In the current version of the function, there will be an error if you forget to specify the value of AT.

Test:

```
GC(dna) Error in GC(dna) : argument "AT" is missing,  
with no default
```



# Giving a default value to an argument

In the current version of the function, there will be an error if you forget to specify the value of AT.

Test:

```
GC(dna) Error in GC(dna) : argument "AT" is missing,  
with no default
```

We should give the value FALSE per default to AT and it will be changed only if the user specifies AT = TRUE.

# Giving a default value to an argument

In the current version of the function, there will be an error if you forget to specify the value of AT.

Test:

```
GC(dna) Error in GC(dna) : argument "AT" is missing,  
with no default
```

We should give the value FALSE per default to AT and it will be changed only if the user specifies AT = TRUE.

```
GC <- function(dna, AT = FALSE ) etc
```

# Giving a default value to an argument

In the current version of the function, there will be an error if you forget to specify the value of AT.

Test:

```
GC(dna) Error in GC(dna) : argument "AT" is missing,  
with no default
```

We should give the value FALSE per default to AT and it will be changed only if the user specifies AT = TRUE.

```
GC <- function(dna, AT = FALSE ) etc
```

Test:

```
GC(dna) [1] 0.3
```

# Giving a default value to an argument

In the current version of the function, there will be an error if you forget to specify the value of AT.

Test:

```
GC(dna) Error in GC(dna) : argument "AT" is missing,  
with no default
```

We should give the value FALSE per default to AT and it will be changed only if the user specifies AT = TRUE.

```
GC <- function(dna, AT = FALSE ) etc
```

Test:

```
GC(dna) [1] 0.3
```

```
GC(dna, AT=TRUE) [1] 0.7
```

# Returning several values

To do so use a vector or a list.

# Returning several values

To do so use a vector or a list.

```
ci.norm <- function(x,conf=0.95)
{
  q <- qnorm(1-(1-conf)/2)
  return(
list(lower=mean(x)-q*se(x),upper=mean(x)+q*se(x)))
}
```

# Returning several values

To do so use a vector or a list.

```
ci.norm <- function(x,conf=0.95)
{
  q <- qnorm(1-(1-conf)/2)
  return(
list(lower=mean(x)-q*se(x),upper=mean(x)+q*se(x)))
}
```

```
ci.norm(rnorm(100))
$lower [1] -0.1499551
$upper [1] 0.2754680
```

```
ci.norm(rnorm(100,conf=0.99))
$lower [1] -0.1673693
$upper [1] 0.2443276
```

# Contents

- 1 Writing your own functions
- 2 **sapply() and tapply()**
- 3 How to avoid slow R code



# apply() and tapply()

You use `apply()` and its derivatives to apply the same function to each element of an object.

```
v <- 1:4
```

```
sapply(v,factorial)
```

```
# returns a vector, lapply() would return a list
```

```
[1] 1 2 6 24
```

# supply() and tapply()

You use `apply()` and its derivatives to apply the same function to each element of an object.

```
v <- 1:4
```

```
sapply(v,factorial)
```

```
# returns a vector, lapply() would return a list
```

```
[1] 1 2 6 24
```

`tapply()` is used for data frames.

# supply() and tapply()

You use `apply()` and its derivatives to apply the same function to each element of an object.

```
v <- 1:4  
supply(v,factorial)  
# returns a vector, lapply() would return a list  
[1] 1 2 6 24
```

`tapply()` is used for data frames.

Example: data frame containing lifespan for people from 3 classes of weight. You want the mean lifespan for each class.

```
tapply(lifespan,weightcls,mean)
```

```
1 2 3  
69 61 53
```

# Contents

- 1 Writing your own functions
- 2 `sapply()` and `tapply()`
- 3 How to avoid slow R code

# How to avoid slow R code

- R has to interpret your commands each time you run a script and it takes time to determine the type of your variables.
- So avoid using loops and calling functions again and again if possible
- When you use loops, avoid increasing the size of an object (vector ...) at each iteration but rather define it with full size before.
- Think in whole objects such as vectors or lists and apply operations to the whole object instead of looping through all elements.